A single-ancilla syndrome circuit with circuit distance 12 for the gross code

Guancheng Pan[1, *]

[1] *Zhejiang University, Hangzhou, China*

(Dated: October 15, 2026)

The gross code $[[144, 12, 12]]$ is the flagship bivariate-bicycle code, but what a syndrome-extraction circuit actually protects against is its circuit distance d_{circ}, the minimum number of faults that cause an undetected logical error. Strikis, Browne and Beverland proved that no non-interleaved single-ancilla circuit for this code reaches $d_{\mathrm{circ}} = 12$ and asked whether any interleaved circuit can. We give one: an interleaved single-ancilla circuit with $d_{\mathrm{circ}} = 12$ under standard circuit-level depolarizing noise, in both memory bases and for every number of rounds. The construction keeps the extended code of their distance-11 circuit, and hence its distance 11, fixed and changes only the order and timing of CNOTs so that none of the minimum-weight extended-code logicals becomes an undetected circuit error; the proof combines two short arguments with an exact computation. The price is a period of 18 time steps per round instead of 8; in a bulk Pauli-frame model of lifting, no colour-symmetric schedule of this kind has a shorter period. The answer is structural: with idle noise and practical decoders, the longer period makes this circuit's logical error rate higher than that of the existing circuits at every noise level we sampled, and the distance gain pays off only when idle errors are much weaker than gate errors (without idle noise the circuit beats the original one at $p = 4 \times 10^{-3}$; with idle errors at one tenth of the gate error rate it already loses). Combining their bounds with an exhaustive computation of the extended-code distance also confirms their conjecture that the original circuit has $d_{\mathrm{circ}} = 11$.

Bivariate-bicycle (BB) codes [1] are leading candidates for low-overhead quantum memory; IBM's architecture is built on the gross code, with parameters $[[144, 12, 12]]$ [1, 2]. A code is only as good as the circuit that measures its checks: faults inside the circuit spread to several qubits, and the relevant figure of merit is the circuit distance d_{circ}, the minimum number of circuit faults that flip a logical observable without triggering any detector. At low physical error rates the logical error rate is governed by d_{circ}, not by the code distance d.

No circuit for the gross code had been shown to reach $d_{\mathrm{circ}} = d = 12$. The IBM circuit, with 8 time steps per round, has $d_{\mathrm{circ}} \leq 10$ [1, 2]. Strikis, Browne and Beverland (SBB) [3] analysed single-ancilla circuits, in which each check is measured by one dedicated ancilla coupled by CNOTs. They append to the check matrix the "residual" errors that an ancilla fault spreads onto the data, obtaining an extended code of distance d_{ext}, and show that $d_{\mathrm{circ}} \geq d_{\mathrm{ext}}$ for every such circuit (their Prop. 3), while $d_{\mathrm{circ}} \leq d_{\mathrm{ext}}$ holds when the circuit is *non-interleaved* in their sense, i.e. when all X-type CNOTs precede all overlapping Z-type CNOTs or vice versa (their Sec. II A and Cor. 1). Since every CNOT order gives $d_{\mathrm{ext}} \leq 11$ for the gross code, non-interleaved circuits are capped at 11 (their Prop. 4); their depth-8 circuit (their Fig. 12, hereafter "Fig. 12") was conjectured to reach it. They left open "whether any interleaved circuit can achieve circuit distance 12" [3].

We answer this question in the affirmative. The circuit of Table I is single-ancilla, CNOT-based and interleaved, and has $d_{\mathrm{circ}} = 12$ for every number of rounds (Theorem 1). The idea (Fig. 1) is to exploit the gap between Prop. 3 and Cor. 1: we keep the extended code of Fig. 12, with $d_{\mathrm{ext}} = 11$, and choose CNOT orders and times such that none of its minimum-weight logicals can be realised by eleven undetected circuit faults. The cost is time: the period is 18 steps per round instead of 8. Our answer is therefore an existence result about circuit structure; under circuit-level noise with idling and with practical decoders, the longer period outweighs the gain in distance at the noise levels we sampled (Discussion). Together with SBB's Prop. 3 and Cor. 1, our computation of d_{ext} confirms their conjecture that Fig. 12 has $d_{\mathrm{circ}} = 11$.

Exact circuit distances are known for hypergraph-product codes [4], for the smallest BB circuit from exact solvers [5], and as certificates for other qLDPC constructions [6, 7]; for larger BB circuits reported values are typically upper bounds from randomized search. Shaw and Terhal's morphing circuit for the $[[288, 12, 18]]$ code uses the same number of qubits, has no dedicated ancillas, and at the end of each cycle acts on a $[[144, 12, 12]]$ code that they call "a second 'gross' code"; its circuit distance is reported as an upper bound, ≤ 12 [8]. Ref. [10] (App. C) describes a morphing idling circuit for the gross code that is "believed to have circuit distance 12", without giving the circuit; presumably it is this one. We note that this end-cycle code is not equivalent to the gross code under qubit permutations and local Cliffords, since its minimum stabiliser weight is 9 rather than 6. SBB also show that the distance of an interleaved circuit is lower-bounded by that of its non-interleaved counterpart [3]; Ref. [11] proposed using this bound to screen interleaved schedules, leaving the search to future work. For the gross code that bound gives at most 11, so our route through a fixed extended code is different. Check-dependent CNOT orders, by check type [12] or by colour class [3, 13], are an established tool.

Setting. We consider memory experiments in the Z or X basis: data initialisation, R rounds of syndrome extraction and transversal readout. Each check has one ancilla that is reset, coupled by six CNOTs and measured once per round; rounds of different checks may overlap in time (pipelining). Detectors compare consecutive measurements of a check, plus the first round and the final readout. Faults are the single-operation Pauli components of standard circuit-level depolarizing noise (SD6 [14]: two-qubit depolarization after CNOTs, flips after resets and before measurements, single-qubit depolarization on idle qubits). Faults with identical detector and observable effects are merged into one detector-error-model (DEM) mechanism [15], and d_{circ} counts distinct mechanisms: it is the minimum number of DEM mechanisms that flip a logical observable while triggering no detector. Lower bounds hold for any sub-model (e.g. without idle noise).

For a check with CNOT order $(q_1, \ldots, q_6)$, an ancilla fault after the k-th CNOT deposits the residual $\{q_{k+1}, \ldots, q_6\}$ on the data. The extended code of a basis appends all residual columns (Hs, Ls) of the relevant checks to the data columns (H: check matrix, L: logical operators). Every single fault, propagated to the readout, has a data image that is zero or one extended-code column (SBB, Table I and the proof of their Prop. 3; proved for our template in SM S5). For Fig. 12, $d_{\mathrm{ext}} = 11$ in both bases, with exactly 4224 weight-11 extended logicals per basis, all realisable in Fig. 12 by eleven undetected faults.

Key idea. Interleaving decouples d_{circ} from d_{ext} in one direction only: $d_{\mathrm{circ}} \geq d_{\mathrm{ext}}$ still holds, but d_{circ} can exceed it. To use this, the extended code must stay at distance 11, and the timing must prevent every weight-11 extended logical from lifting. We say that an extended logical *lifts* if one can choose, for each of its columns, a circuit fault whose propagated image is that column, such that the chosen faults together trigger no detector; a lift is then an undetected logical error of the same weight. The first requirement is met by an invariance. Up to the check itself, a residual of length 5 equals $\{q_1\}$, of length 4 equals $\{q_1, q_2\}$, of length 3 equals the complement of $\{q_1, q_2, q_3\}$, and of length 2 is $\{q_5, q_6\}$. The extended code therefore depends only on the three unordered sets $\{q_1, q_2\}$, $\{q_1, q_2, q_3\}$, $\{q_5, q_6\}$ of each check, which are preserved by reversal (rev) and by the swaps s_{12}, s_{56} of positions 1,2 and 5,6. All eight orders generated by these operations share the extended code of Fig. 12, its distance 11 and its list of 4224 weight-11 logicals (Fig. 1a).

The second requirement is a statement about time. Reversing an order keeps the extended code but changes which physical fault realises each residual column and when its spread reaches the neighbouring checks. Consider a weight-11 extended logical made of ten data errors and one two-qubit hook column $\{q_5, q_6\}$ of an X check (Fig. 1b). In Fig. 12 the column is produced by an ancilla fault after the fourth CNOT, which lands on the data after the neighbouring Z checks have read those qubits; every data error can then be placed after its reads, and no detector fires. With the order reversed, the same column can only come from a fault between the second and third CNOT, whose spread over four qubits (equal to the hook up to a stabiliser) reaches one neighbouring check after its readout and another before it. The checks along a cycle through the hook and three data qubits each see two errors of the logical and stay silent only if both fall on the same side of their readout; the early and late constraints propagate around the cycle and leave one check with odd parity, so a detector fires: timing turns the potential weight-11 logical into a parity cycle that cannot be satisfied. For all 360 logicals of this "ten data plus one hook" form, the lift fails at $R = 1$, and every irreducible conflicting set of detectors (one whose constraints alone forbid the lift, but no proper subset does) has between 4 and 10 detectors; the smallest such set has 4, 5 or 9 detectors. Circuit design becomes a finite problem: find orders in the class and CNOT times for which none of the 4224 known logicals lifts. With the Fig. 12 orders themselves this is impossible in every family of retimings we could decide (SM S9).

Construction. We searched schedules that are symmetric under the 24 translations preserving the colouring, choosing for each of the six (type, colour) classes one of the eight orders and its CNOT times with a SAT solver; constraints enforce legality (no data qubit in two CNOTs at once), determinism of all detectors, and blocking of every weight-11 extended logical in a bulk Pauli-frame model (SM S10). Candidates were then certified exactly. Table I and Fig. 2 give the resulting circuit. It measures the gross code itself, is legal and deterministic, is pipelined (every check's window crosses a period boundary) and interleaved (96 data qubits switch between X- and Z-type CNOTs four times per period, 48 twice); every check's order differs from that of Fig. 12. A second period-18 circuit and a period-21 circuit in which all ancillas are reset and measured together have the same distance (SM S2).

Theorem 1 *Under SD6 noise, the circuit of Table I has* $d_{\mathrm{circ}} = 12$ *in both memory bases for every number of rounds* $R \geq 1$.

The proof has three steps, two of them pen-and-paper and one computer-assisted.

Lemma 1 (Reduction) *In any circuit of our template, an undetected fault set of size* w *that flips a logical observable yields an extended-code logical of weight at most* w*, formed by its nonzero fault images.*

Proof. Add one observable per row of H and L, evaluated on the final data readout. The observable of row H_i equals the sum of all detectors of check i, since these

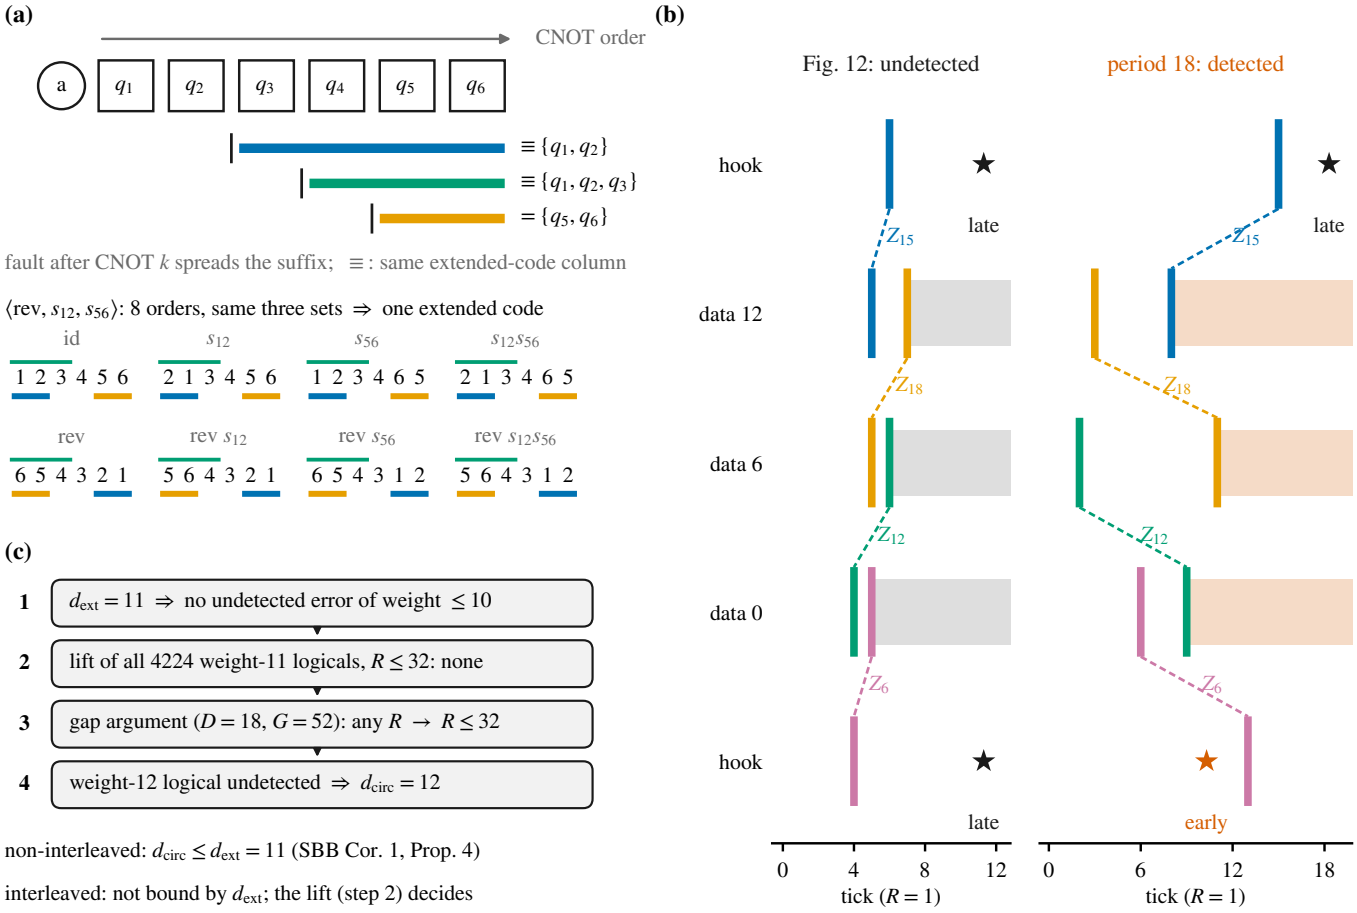


FIG. 1. The idea. (a) An ancilla fault after the k-th CNOT of a check spreads the residual $\{q_{k+1}, \ldots, q_6\}$ onto the data. Up to the check itself, the residuals of length 4, 3 and 2 give the same extended-code columns as $\{q_1, q_2\}$, $\{q_1, q_2, q_3\}$ and $\{q_5, q_6\}$, and those of length 5 and 1 act as single data qubits. The eight orders generated by reversal and the swaps s_{12}, s_{56} (top row: id, s_{12}, s_{56}, $s_{12}s_{56}$; bottom row: the same after reversal, e.g. "rev s_{12}" = reverse, then swap) keep these three sets and therefore share one extended code. (b) One weight-11 extended logical in a Z-basis memory with $R = 1$: ten data errors and the hook column $\{9, 25\}$ of an X check. Rows: the hook, data qubits 12, 6 and 0, and the hook again, closing the cycle; bars: the reads of the Z checks Z_{15}, Z_{18}, Z_{12}, Z_6, joined pairwise for each check; stars: when the hook error reaches the data, labelled late or early. Each of these checks sees exactly two errors of the logical and is silent only if both lie on the same side of its read. Left, Fig. 12: the hook lands on qubit 9 at tick 11, after its reads at ticks 6 and 4, and all data errors can be placed after their reads (grey); no detector fires. Right, period-18 circuit: the order of this check is reversed, so the residual after its fourth CNOT is no longer $\{9, 25\}$; the same extended-code column is realised only by an ancilla fault between the second and third CNOT, which spreads to the four qubits $\{8, 91, 82, 85\}$, equal to $\{9, 25\}$ up to the check itself. This spread is seen by Z_{15} on qubit 85 at tick 18, after its read at 15 (late), and by Z_6 on qubit 82 at tick 10, before its read at 13 (early). Being late for Z_{15} forces the data errors late for Z_{18}, Z_{12} and Z_6 (red), which leaves odd parity at Z_6, so a detector fires. (c) Logic of the proof: $d_{\text{ext}} = 11$ excludes weight ≤ 10 (Lemma 1); none of the 4224 weight-11 extended logicals lifts for $R \leq 32$ (Lemma 2); the gap argument reduces every R to $R \leq 32$ (Lemma 3); a weight-12 witness gives the upper bound. Non-interleaved circuits obey $d_{\text{circ}} \leq d_{\text{ext}} = 11$; for interleaved circuits the lift decides.

telescope from the first measurement to the final readout. For an undetected fault set the H parts of the images therefore sum to zero, while the L parts sum to the logical effect. Each image is zero or one column, a property noted by SBB (their Table I) that we prove for our template, interleaved and pipelined, in SM S5; pairs of equal columns cancel. □ This uses neither Prop. 3 of SBB nor non-interleaving; with $d_{\text{ext}} = 11$ it excludes weight ≤ 10 for every R.

Lemma 2 (Finite rounds) *For $R = 1, \ldots, 32$ and both bases, none of the 4224 weight-11 extended logicals is the image of eleven undetected faults.*

We verified this by an exact computation: for each logical, a CP-SAT model chooses one DEM mechanism per column subject to every detector sum vanishing; all $2 \times 32 \times 4224$ instances are infeasible. As a positive control, the same program finds a realisation for all 4224 logicals in Fig. 12. The computation was carried out

TABLE I. The period-18 circuit. Checks are grouped by type and colour $(x+y)$ mod 3 of their position (x, y) on the 12×6 torus; all 24 checks of a class share reset step, CNOT order (labels $A = x^3 + y + y^2 \mapsto A_1A_2A_3$, $B = y^3 + x + x^2 \mapsto B_1B_2B_3$), CNOT steps and measurement step, all modulo 18. A window that passes step 17 continues in the next period. Last column: map from the Fig. 12 order of the class.

class	reset	CNOT order	CNOT steps	meas.	map
X_0	17	$A_1A_2A_3B_3B_1B_2$	0, 1, 2, 8, 9, 15	16	rev s_{56}
X_1	15	$A_1A_2A_3B_1B_3B_2$	16, 3, 4, 5, 10, 13	14	rev
X_2	17	$A_3A_1A_2B_3B_1B_2$	0, 1, 5, 6, 7, 15	16	rev s_{56}
Z_0	17	$B_2B_1B_3A_2A_1A_3$	0, 7, 8, 9, 13, 14	15	rev s_{56}
Z_1	2	$B_2B_3B_1A_3A_2A_1$	3, 4, 9, 10, 11, 17	0	rev $s_{12}s_{56}$
Z_2	16	$B_2B_1B_3A_3A_2A_1$	17, 5, 6, 12, 13, 14	15	rev s_{56}

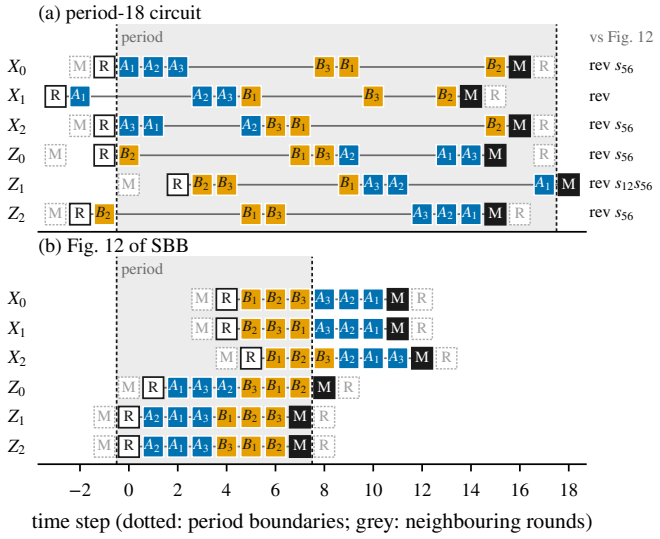


FIG. 2. One period of (a) the circuit of Table I and (b) Fig. 12 of SBB [3], on the same time scale. Rows: the six classes; boxes: CNOTs labelled by the term they implement; R, M: ancilla reset and measurement. Windows crossing the period boundary (dotted) continue into the next period. (b) has period 8 and is not interleaved: on every data qubit the X- and Z-type CNOTs of a round form two contiguous blocks.

with two independent implementations, and the inputs ($d_{\text{ext}} = 11$ and the list of 4224 logicals) with two independent exhaustive searches (SM S4, S6). The infeasible lift instances of this finite-R computation were also encoded as CNF and combined into batches by a selector disjunction, so that unsatisfiability of a batch implies that of each member; each batch was refuted with a DRAT proof verified with drat-trim, whose LRAT core was verified with cake_lpr (SM S6).

Lemma 3 (All rounds) *If a weight-11 undetected logical existed for some R, one would exist for some $R' \leq 32$.*

Proof sketch. After the last fault of a fault set at layer a, the data frame is constant from layer $a + D - 1$ on, $D = 18$, because only ancillas still in flight can change it. If an undetected logical had no fault in a window of $G = 2D + \Delta - 1 = 52$ layers, where $\Delta = 17$ is the longest reset-to-measurement span of a check, no detector could respond to both the faults before and after the window; both parts would be undetected, their observables would add to a nonzero value, and one part would have weight ≤ 10, contradicting Lemma 1. So a weight-11 logical spans at most 10×51 layers. The noisy circuit is periodic: inserting an 18-layer block at any layer x with $4 \leq x \leq 18R$ (layers numbered from 0) turns the R-round circuit into the $(R+1)$-round one. Deleting fault-free blocks before and after the logical leaves every detector it touches unchanged, which maps it into a circuit with at most 32 rounds (SM S7). □

Lemmas 1–3 give $d_{\text{circ}} \geq 12$ for every R. For the upper bound, a weight-12 logical operator of the code, applied as single-qubit flips after data initialisation, triggers no detector and flips the logical observable (SM S8). This proves Theorem 1.

Corollary 1 *The Fig. 12 circuit of SBB has $d_{\text{circ}} = 11$ in both bases for every R, as they conjectured.*

This is SBB's Prop. 3 ($d_{\text{circ}} \geq d_{\text{ext}}$) and Cor. 1 ($d_{\text{circ}} \leq d_{\text{ext}}$ for non-interleaved circuits) combined with our value $d_{\text{ext}} = 11$, obtained by a complete symmetric search over the extended code and reproduced with an independently built instance and a second symmetry group. Retiming Fig. 12 without changing its orders does not help: among schedules of period ≤ 8 there are four conflict-free ones up to translation, all with weight-11 undetected errors, and in the bulk, every legal circuit keeping the orders has such an error in each longer-period family of Table S3 (SM S9).

Cost and optimality. The circuit uses the qubits of the IBM and SBB circuits, one ancilla per check, but 18 time steps per round instead of 8, i.e. 2.25 times as many, with more idle locations. Whether a shorter period is possible we can answer only within a model. In a bulk Pauli-frame model of lifting (SM S9), a data error is seen by a check at the first readout after it, and a weight-11 extended logical is *model-liftable* if its errors can be

timed so that every check sees an even number of them between consecutive readouts. For logicals of ten data errors and one residual of length two, a model lift yields a genuine weight-11 undetected error once the number of rounds is large enough (a one-dimensional Helly argument, SM S9); for other logicals we confirmed this on the exact detector error model in all 991 sampled cases, including logicals with several residual columns. This is the direction a non-existence claim needs. Among schedules symmetric under the 24 colour-preserving translations with orders in the class of Fig. 12, the model admits no schedule with period $P \leq 17$ in which no weight-11 logical is model-liftable, and none with $P \leq 20$ when all ancillas share one reset and measurement frame; within this model the minimal periods are therefore 18 and 21. Because the model-to-circuit step is proven only for logicals with one length-2 residual, we regard this as strong evidence rather than a theorem. The class of orders is not arbitrary either: among all 24 colourings by cosets of translation subgroups of index 2, 3, 4, 6, 8 or 9, with one order type per colour, only the class of Fig. 12 has $d_{\text{ext}} \geq 11$ in both bases, an exact statement. With less symmetry the minimal period is open; in the same model, invariance under a translation subgroup of order 12 already excludes $P \leq 11$, and any non-trivial subgroup excludes $P = 8$.

Discussion. The obstruction found by SBB is specific to non-interleaved circuits, not to the single-ancilla model: the gross code admits a single-ancilla circuit that preserves its full distance. The method transfers. Fixing the extended code and searching only over orders within its class and over timing turns an open-ended circuit search into a lift problem against a finite list of minimum-weight extended logicals, which can be decided exactly and certified; the same approach applies to other single-ancilla circuits whose distance is limited by Cor. 1 of SBB.

Practical performance. We compared the period-18 circuit with Fig. 12 and the IBM circuit using the same template, the same SD6 noise (idle depolarization on every idle qubit in every layer), $R = 12$, and identical BP-OSD and Relay-BP decoders with exact priors. Under this model the period-18 circuit is worse at every sampled $p \geq 10^{-3}$: it is above its pseudo-threshold for $p \geq 3 \times 10^{-3}$, about 1.5×10^3 times worse than the IBM circuit at $p = 2 \times 10^{-3}$, and at least 3 times worse at 10^{-3} (a direct-sampling bound). The reason is fault mass: the longer period has 3168 idle locations per round instead of 288, eleven times as many, and 2.8 times the total fault probability per round; at matched fault mass the circuits perform alike. Raising d_{circ} from 11 to 12 does not change the leading power of an ideal decoder ($\lceil d_{\text{circ}}/2 \rceil = 6$ for both), and for the enumerated families of weight-12 logicals the leading-order coefficient of the period-18 circuit is about 26 times that of Fig. 12. $d_{\text{circ}} = 12$ is thus a structural property of the circuit and not a practical improvement under this noise model.

Role of idle noise. To isolate it we rescale only the idle channel of SD6, $p_{\text{idle}} = \lambda p$, leaving all other components unchanged ($\lambda = 1$ reproduces SD6 exactly; $\lambda = 0$ coincides with circuit-level depolarizing noise without idle errors). At $\lambda = 0$ the period-18 circuit has a lower per-round logical error rate than Fig. 12 at $p = 4 \times 10^{-3}$ (ratio 0.37 [0.28, 0.49] with BP-OSD, independently reproduced; 0.25 [0.11, 0.60] with Relay-BP) and is statistically indistinguishable from the IBM circuit at $p = 3$–4×10^{-3}. Already at $\lambda = 0.1$ it is worse than Fig. 12 by a factor 2.3–3 per round for $2 \times 10^{-3} \leq p \leq 4 \times 10^{-3}$ (comparable per unit-duration layer), and at $\lambda = 0.3$ by 15–70 times. A leading-order (ideal-decoder) estimate restricted to the enumerated weight-11 and weight-12 logical families places the per-round break-even with Fig. 12 near $\lambda \approx 0.3$; because one type of weight-12 logical is not enumerated for either circuit, this value is indicative only (an unknown factor of 3 in their relative weight moves it to 0.2–0.5), and it concerns the asymptotic regime $p \lesssim 10^{-4}$, not the sampled range. Per-layer rates assume unit-duration layers; with realistic measurement and reset durations the advantage the per-layer comparison gives the longer period largely disappears. A variant with SI1000 coefficients and native CX (CX, RX and MX treated as native) gives the same qualitative picture (SM S11).

Our results do not cover correlated faults, leakage or CZ-native compilations. The main open question is whether $d_{\text{circ}} = 12$ can be reached at period 8, or at any period below 18, by less symmetric schedules. Since the distance gain pays off only when idle errors are much weaker than gate errors (at $\lambda = 0$ the circuit beats Fig. 12 at $p = 4 \times 10^{-3}$, at $\lambda = 0.1$ it already loses), a shorter period is the main open practical question.

Data availability. The circuit specifications, the lists of 4224 weight-11 extended logicals for each basis, the witnesses, the certificates (clause sets of the unsatisfiable instances) and the reproduction scripts are available at [Zenodo DOI — PI to decide].

AI use. This research was carried out with AI agents (Claude, Anthropic). The agents proposed hypotheses, wrote and ran the search and verification code, worked out the constructions and proofs, and independently rechecked each claim reported here, using separately written code for every computation. The author directed the project, made all decisions at its decision points, and takes full responsibility for the content. In line with APS policy, no AI system is listed as an author.

* 3250101086@zju.edu.cn

[1] S. Bravyi, A. W. Cross, J. M. Gambetta, D. Maslov, P. Rall, and T. J. Yoder, High-threshold and low-

overhead fault-tolerant quantum memory, Nature **627**, 778 (2024), arXiv:2308.07915.

[2] T. J. Yoder, E. Schoute, P. Rall, E. Pritchett, J. M. Gambetta, A. W. Cross, M. Carroll, and M. E. Beverland, Tour de gross: A modular quantum computer based on bivariate bicycle codes, arXiv:2506.03094 (2025).

[3] A. Strikis, D. E. Browne, and M. E. Beverland, High-performance syndrome extraction circuits for quantum codes, arXiv:2603.05481 (v2, 2026).

[4] A. G. Manes and J. Claes, Distance-preserving stabilizer measurements in hypergraph product codes, Quantum **9**, 1618 (2025), arXiv:2308.15520.

[5] M. Webster, A. Jacob, and O. Higgott, Distance-finding algorithms for quantum codes and circuits, arXiv:2603.22532 (v2, 2026).

[6] Z. Liang, B. Gu, Y.-A. Chen, J. Eisert, and Z. Wang, Ultra-high-distance quantum memories from amplified qLDPC codes, arXiv:2609.37231 (2026).

[7] Z. Sun and Z. Cai, Layer codes as quantum memories: syndrome extraction, thresholds and idle robustness, arXiv:2609.39522 (v2, 2026).

[8] M. H. Shaw and B. M. Terhal, Lowering connectivity requirements for bivariate bicycle codes using morphing circuits, Phys. Rev. Lett. **134**, 090602 (2025), arXiv:2407.16336.

[9] M. H. Shaw and B. M. Terhal, Optimising quantum error correction using morphing circuits, arXiv:2604.09797 (v2, 2026).

[10] A. Wills, M. E. Beverland, L. S. Bishop, J. M. Gambetta, P. Rall, V. Siddhu, and A. W. Cross, Concatenating algebraic codes over high-rate quantum LDPC codes, arXiv:2605.21898 (v2, 2026).

[11] M. T. P. Nguyen, M. Rimbach-Russ, and S. Bosco, Disassembling qLDPC codes for depth-optimal parity-check circuits, arXiv:2608.19917 (v2, 2026).

[12] Y. Tomita and K. M. Svore, Low-distance surface codes under realistic quantum noise, Phys. Rev. A **90**, 062320 (2014), arXiv:1404.3747.

[13] G. Kishony and A. Fowler, Color code off-the-hook: avoiding hook errors with a single auxiliary per plaquette, arXiv:2603.28852 (2026).

[14] C. Gidney, M. Newman, A. Fowler, and M. Broughton, A fault-tolerant honeycomb memory, Quantum **5**, 605 (2021), arXiv:2108.10457.

[15] C. Gidney, Stim: a fast stabilizer circuit simulator, Quantum **5**, 497 (2021), arXiv:2103.02202.

Supplemental Material

The sections follow the order of the main text. S1–S2 fix notation and the circuits; S3–S8 contain the proof of Theorem 1; S9 concerns Fig. 12 of SBB; S10 the search and the optimality statements; S11 the sensitivity of the logical error rate to idle noise; S12 the verification and the data.

S1. NOTATION AND NOISE MODEL

Code. The gross code [1] has $\ell = 12$, $m = 6$, $A = x^3 + y + y^2$, $B = y^3 + x + x^2$, $H_X = [A \,|\, B]$, $H_Z = [B^T \,|\, A^T]$; qubits and checks are indexed by $i = 6a + b$ for the monomial $x^a y^b$, with left and right blocks of 72 qubits each; $n = 144$, $k = 12$, $d = 12$. Term labels: $A_1 = x^3$, $A_2 = y$, $A_3 = y^2$, $B_1 = y^3$, $B_2 = x$, $B_3 = x^2$. The X check at (x, y) acts on left qubit $(x, y) + A_k$ and right qubit $(x, y) + B_k$; the Z check at (x, y) on left qubit $(x, y) - B_k$ and right qubit $(x, y) - A_k$. The colour of a check is $(x + y)$ mod 3; the 24 translations (a, b) with $a + b \equiv 0 \pmod 3$ preserve colours.

Circuit template. Each check has one ancilla. In every round it is reset, applies its six CNOTs in its order (ancilla as control for X checks, as target for Z checks) and is measured. A schedule assigns each check a reset step, six CNOT steps and a measurement step modulo the period P; windows may wrap around the period boundary, so that rounds of different checks overlap. A Z-basis memory initialises the data in $|0\rangle$, runs R rounds and reads all data in the Z basis; detectors are the first measurement of each Z check, the comparison of consecutive measurements of each Z check, and the comparison of the last measurement with the parity of the final readout. The X basis is analogous.

Noise and d_{circ}. SD6 [14]: after every CNOT a two-qubit depolarizing channel; after every reset a flip; before every measurement a flip; on every qubit idle in a time step a single-qubit depolarizing channel (also during reset and measurement layers). A fault is one Pauli component of one channel. Faults with identical detector and observable signatures are one DEM mechanism, and d_{circ} counts mechanisms. For lower bounds we use the CSS relaxation: in a Z-basis memory only Z-check detectors and the X part of each fault are kept; this can only lower the minimum. Upper bounds are checked on the full circuit. Lower bounds hold for any sub-model of SD6.

S2. THE CIRCUITS

Table I specifies circuit A completely: each of the 144 checks takes the reset step, order, CNOT steps and measurement step of its class. Its resets occur at steps $\{2, 15, 16, 17\}$ and its measurements at $\{0, 14, 15, 16\}$. We verified for circuits A, B and C that every check applies exactly six CNOTs to its support at distinct steps inside its window, that no data qubit takes part in two CNOTs at the same step, that the noiseless circuit fires no detector, and that all detectors are deterministic, for several R in both bases. Interleaving: in circuit A, over one period, 96 data qubits switch four times and 48 twice (cyclically) between X- and Z-type CNOTs.

Circuit B (Table S1) also has period 18; 24 of its checks keep the Fig. 12 order. Circuit C (Table S2) has period 21 and a common frame: all ancillas are reset at step 0 and measured at step 20, so rounds are separated in time. Both have $d_{\mathrm{circ}} = 12$ for every R in both bases (S6, S7).

For reference, the Fig. 12 orders in the same labelling are X_0: $B_1B_2B_3A_3A_2A_1$, X_1: $B_2B_3B_1A_3A_2A_1$, X_2: $B_1B_2B_3A_2A_1A_3$, Z_0: $A_1A_3A_2B_3B_1B_2$, Z_1: $A_2A_1A_3B_1B_2B_3$, Z_2: $A_2A_1A_3B_3B_1B_2$ (period 8, pipelined, non-interleaved).

TABLE S1. Circuit B (period 18). Columns as in Table I.

class	reset	CNOT order	CNOT steps	meas.	map
X_0	0	$B_2B_1B_3A_3A_2A_1$	1, 7, 8, 14, 15, 16	17	s_{12}
X_1	2	$B_2B_3B_1A_3A_2A_1$	3, 6, 11, 12, 13, 0	1	id
X_2	0	$B_2B_1B_3A_2A_1A_3$	1, 9, 10, 11, 15, 16	17	s_{12}
Z_0	1	$A_3A_1A_2B_3B_1B_2$	2, 3, 7, 8, 9, 17	0	s_{12}
Z_1	17	$A_1A_2A_3B_1B_3B_2$	0, 5, 6, 7, 12, 15	16	$s_{12}s_{56}$
Z_2	1	$A_1A_2A_3B_3B_1B_2$	2, 3, 4, 10, 11, 17	0	s_{12}

TABLE S2. Circuit C (period 21, common frame: reset at 0, measurement at 20).

class	CNOT order	CNOT steps	map
X_0	$B_2B_1B_3A_3A_2A_1$	2, 7, 9, 15, 16, 17	s_{12}
X_1	$B_2B_3B_1A_3A_2A_1$	1, 8, 12, 13, 14, 17	id
X_2	$B_2B_1B_3A_2A_1A_3$	1, 10, 11, 12, 16, 18	s_{12}
Z_0	$A_3A_1A_2B_3B_1B_2$	3, 4, 8, 9, 10, 17	s_{12}
Z_1	$A_1A_2A_3B_1B_3B_2$	1, 6, 7, 8, 13, 14	$s_{12}s_{56}$
Z_2	$A_1A_2A_3B_3B_1B_2$	2, 4, 5, 11, 12, 19	s_{12}

S3. INVARIANCE OF THE EXTENDED CODE

Consider a Z-basis memory; only X errors matter. An X error on the ancilla of an X check after its k-th CNOT propagates to the suffix $\{q_{k+1}, \dots, q_6\}$; Z-check CNOTs (data controls, ancilla target) do not change the X frame on the data. Each fault's final data X frame is therefore zero, a single data qubit, or a proper suffix of an X-check order, and its column is $(H_Z s, L_Z s)$. Multiplying by the check itself, which has zero column, a suffix of length 5 is equivalent to $\{q_1\}$ (a data column), of length 4 to $\{q_1, q_2\}$, of length 3 to the complement of $\{q_1, q_2, q_3\}$, and a suffix of length 2 is $\{q_5, q_6\}$. Hence the column set is determined by the three sets $\{q_1, q_2\}$, $\{q_1, q_2, q_3\}$, $\{q_5, q_6\}$ of every check. Reversal maps these to $\{q_5, q_6\}$, the last three and $\{q_1, q_2\}$, which produce the same columns, and s_{12}, s_{56} leave them unchanged. The X basis is symmetric. For circuits A, B and C we also computed the extended codes directly: both bases have 360 merged columns, and the column sets coincide with those of Fig. 12.

S4. THE EXTENDED CODE: $d_{\text{ext}} = 11$ AND THE 4224 WEIGHT-11 LOGICALS

The extended code of Fig. 12 has 360 merged columns per basis (144 of check-weight 3, 144 of weight 6, 72 of weight 9). The 24 colour-preserving translations map each check's order to the order of its image check and preserve the logical classes (checked by solving $L \circ T = ML + NH$ with $\operatorname{rank} M = k$); the columns form 15 orbits.

Search. To decide whether $Hx = 0$, $Lx \neq 0$ has a solution with $|x| \leq W$ we use a complete depth-first search: a minimum-weight solution may be translated so that it contains the representative of its lowest orbit; the search branches on the columns covering an unsatisfied check row, excluding earlier siblings, so that the branches partition the search space; it prunes when the depth plus $\lceil |s|/w_{\max} \rceil$ exceeds W (s the current syndrome, $w_{\max}$ the largest column weight) and closes the last step by table look-up. The program was validated against unpruned meet-in-the-middle enumeration on smaller instances and by checking that split runs add up to unsplit runs.

Results. At $W = 10$ there is no solution in either basis ($\approx 2 \times 10^{10}$ nodes per basis), on two independently constructed instances, and under a second symmetry group of order 8 (45 orbits, $\approx 5 \times 10^{10}$ nodes per basis); weight 11 is attained. Hence $d_{\text{ext}} = 11$. At $W = 11$, two different search programs enumerate the same 4224 extended logicals per basis, all in free orbits of the 24-element group (176 orbits); 360 of them consist of ten data columns and one residual of length 2. Both $d_{\text{ext}} = 11$ and the completeness of the list of 4224 weight-11 logicals rest on two independent exhaustive searches; we did not obtain a checkable proof certificate for them (a direct CNF encoding is beyond CDCL solvers already at weight 6).

S5. PROOF OF LEMMA 1

We work in a Z-basis memory; the X basis follows by exchanging the roles of X and Z and of the two check types. In a Z-basis memory only X components matter. For a fault we track its *data X frame*: the set s of data qubits carrying an X error after the fault has been propagated through the rest of the circuit to the final data readout. Its *image* is the column $(H_Z s,\, L_Z s)$.

Gates of the template. An X-check ancilla is prepared in $|+\rangle$, is the control of its six CNOTs (data qubits as targets) and is measured in the X basis. A Z-check ancilla is prepared in $|0\rangle$, is the target of its six CNOTs (data qubits as controls) and is measured in the Z basis. Under a CNOT, an X on the control is copied to the target, an X on the target stays there, and Z errors never create X errors. Hence: (a) an X on a data qubit never leaves it (as a target it is not copied; as a control it is copied only onto Z-check ancillas, which are never controls and are measured and

reset without acting back on the data); (b) an X on a Z-check ancilla never reaches the data; (c) an X on an X-check ancilla after its k-th CNOT of an instance is copied onto exactly the targets $q_{k+1}, \ldots, q_6$ of the remaining CNOTs of that instance and is then discarded by the next preparation (the X-basis measurement does not see it). These statements use only the direction of each gate, not the relative timing of different checks, so they hold for interleaved and pipelined schedules alike.

Single faults. We go through the SD6 locations. In each case s is empty, a single data qubit, or a suffix $\{q_k, \ldots, q_6\}$ of one X-check order.

1. Data qubits: an X or Y from idle noise, a flip after data preparation, or a flip before the final readout gives $s = \{q\}$ by (a); a Z gives $s = \emptyset$.

2. Z-check ancillas: preparation and measurement flips, idle errors and any error on the ancilla alone give $s = \emptyset$ by (b).

3. X-check ancillas: the preparation flip and the measurement flip are Z errors in this basis and give $s = \emptyset$; an X or Y from idle noise before the first CNOT gives the full support of the check, after the k-th CNOT the suffix $\{q_{k+1}, \ldots, q_6\}$, after the sixth CNOT $s = \emptyset$, by (c).

4. Two-qubit Pauli $P_a \otimes P_d$ after the k-th CNOT of an X check (ancilla a, data q_k): the X part of P_a gives $\{q_{k+1}, \ldots, q_6\}$ and the X part of P_d gives $\{q_k\}$, so $X_a I$ gives $\{q_{k+1}, \ldots, q_6\}$, $I X_d$ gives $\{q_k\}$, and $X_a X_d$ (or any combination with X parts on both) gives their union $\{q_k, q_{k+1}, \ldots, q_6\}$, the suffix starting at q_k.

5. Two-qubit Pauli after a CNOT of a Z check (data q, ancilla): the data part gives $\{q\}$ or $\emptyset$ by (a), the ancilla part gives $\emptyset$ by (b).

Images. A single data qubit gives a data column of the extended code. A suffix $\{q_k, \ldots, q_6\}$ with $2 \le k \le 6$ is a residual, i.e. one of the appended columns (for $k = 6$ it is the data column of q_6). The full support ($k = 1$) is the check itself, whose image is zero because H_Z and L_Z commute with the X stabilisers. So every single fault has image zero or exactly one extended-code column. Merging faults into DEM mechanisms does not change this, since merged faults have identical effects on all detectors and observables, including the augmented ones introduced below.

Telescoping. Augment the circuit with one observable for each row of H_Z and of L_Z, measured on the final data readout. For every Z check i, the sum of all its detectors (first measurement, all consecutive comparisons, final comparison) telescopes to the parity of row i of H_Z on the final readout, i.e. to the augmented observable H_i. Hence an undetected fault set has images whose H parts sum to zero, and whose L parts sum to its logical effect. Its nonzero images therefore sum to an extended-code logical; cancelling equal pairs leaves a set of distinct columns of weight at most the number of faults. If the set has eleven faults and $d_{\text{ext}} = 11$, the eleven images are nonzero, distinct and form one of the listed weight-11 logicals. The template has no X-check detectors in a Z-basis memory; adding them would only remove undetected fault sets, so the bound also holds with all detectors.

Computed check. For every circuit and every R used in S6 we confirmed both facts on the DEM in both bases: each mechanism's per-check detector parity equals the H part of its image, its logical effect equals the L part, and the set of images is exactly the 360 extended columns.

S6. COMPUTATION FOR LEMMA 2

By S5, a weight-11 undetected logical selects, for each of the eleven columns of one listed extended logical, one DEM mechanism with that image, such that the detector sums vanish. For each of the 4224 logicals per basis, a CP-SAT model has a Boolean variable per mechanism of the eleven image groups, an exactly-one constraint per group, and an XOR constraint per detector. Infeasibility for all 4224 logicals at a given R means that no weight-11 undetected logical exists at that R.

Results: circuit A, $R = 1, \ldots, 32$, both bases: all infeasible. A second, independently written lift program agrees and extends the range to $R = 36$. Circuit B: $R = 1, \ldots, 24$; circuit C: $R = 1, \ldots, 14$ and 16, plus a solver-free meet-in-the-middle check (split the eleven groups in two halves, enumerate all detector sums of each half, intersect) for $R \le 3$; all infeasible. Positive control: on Fig. 12 the same programs find a realisation for all 4224 logicals.

Certificates. For circuits A ($R \le 32$), B ($R \le 24$) and C ($R \le 14$ and 16), both bases and all 4224 logicals, the lift instances were encoded as CNF. Instances were combined into batches by a selector disjunction: the batch formula is the disjunction of its members, each guarded by its own selector variable, with no variables shared between members, so that unsatisfiability of a batch implies that of each member. Each batch was refuted by CaDiCaL; each refutation

is a DRAT proof verified with drat-trim, whose LRAT core was verified with the formally verified checker cake_lpr. As a positive control, every one of the 4224 Fig. 12 instances at $R = 1$ is satisfiable and decodes to a valid lift.

S7. PROOF OF LEMMA 3

Notation: period $D = 18$; check q has start s_q, instance j is reset at $s_q + jD$ and measured at $s_q + jD + \Delta_q$ with $\Delta_q \le D - 1$; $\Delta = \max_q \Delta_q = 17$; $S = \max_q s_q - \min_q s_q = 5$. Faults are attached to layers.

(1) Settling. If all faults of a set $F_<$ lie at layers $\le a$, the data frame they produce is constant from layer $a + D - 1$ on: only an ancilla fault can later change the data, through the remaining CNOTs of its own instance, all at layers $\le a + D - 2$; a freshly reset ancilla never pushes errors onto the data (an X-check ancilla only picks up Z errors from the data, a Z-check ancilla only X errors). Every instance reset after $a + D - 1$ is fault-free and measures $H_c e$. If all faults of $F_>$ lie at layers $\ge b$, measurements before b are unaffected.

(2) Splitting. Let $F = F_< \cup F_>$ be a weight-11 undetected logical with no fault strictly between a and b. A comparison detector (c, j) responds to $F_<$ only if $s(c, j-1) < a + D - 1$ and to $F_>$ only if the measurement of instance j, at $s(c, j-1) + D + \Delta_c \le a + 2D + \Delta - 2$, is at a layer $\ge b$. A first-round detector $(c, 0)$ responds to $F_>$ only if $b \le s_c + \Delta_c \le 21$, so with $a \ge 0$ it responds to both parts only if $b - a \le 21 < G$. A final detector responds to both only if $b - a < D + S + \Delta$. Hence for $b - a \ge G = \max(2D + \Delta - 1,\ D + S + \Delta) = 52$ no detector responds to both parts; each part is undetected, their logical effects add to a nonzero value, and one part has weight ≤ 10, contradicting Lemma 1. Consecutive faults of F are therefore less than G layers apart, and F spans at most $10(G-1) = 510$ layers.

(3) Deletion. The noisy SD6 circuit with R rounds equals the circuit with $R - 1$ rounds with an 18-layer block inserted at layer x, for every $x \in [4, 18(R-1)]$ (checked layer by layer in both bases; the template is exactly periodic). Removing a block $[x, x + D)$ that contains no fault of F removes one measurement of every ancilla; a detector of the shorter circuit that does not straddle the cut equals a detector of the longer one, and one that straddles it equals the sum of the two detectors around the removed measurement. F therefore stays undetected, with the same logical effect, provided the frame entering layer $x + D$ of the longer circuit equals the frame entering layer x of the shorter one; this holds if $x + D \le f_1$ (all zero; f_1 the first fault layer) or $x \ge f_{11} + 2D - 2$ (data frame settled by (1), all ancillas in flight fault-free). *(4) Counting.* Write $f_1 \le f_{11}$ for the first and last fault layers of F; by (2), $f_{11} - f_1 \le 10(G - 1) = 510$. *Front:* while $f_1 \ge 4 + D = 22$, the block $[4, 22)$ is fault-free and can be deleted at $x = 4$, shifting F by -18 layers; afterwards $f_1 \le 21$ and hence $f_{11} \le 21 + 510 = 531$. *Back:* a block can be deleted at $x = f_{11} + 2D - 2$ as long as this x is an admissible insertion point, $x \le 18(R-1)$. When no further deletion is possible, $18(R-1) < f_{11} + 2D - 2 \le 531 + 34 = 565$, so $R - 1 \le 31$ and $R \le 32$. (Using $2D - 1$ in place of $2D - 2$, as a more conservative variant does, gives the same bound; the threshold $2D - 3$ would already suffice.) Thus every weight-11 undetected logical of a circuit with R rounds yields one of a circuit with at most 32 rounds, and with Lemma 2 none exists for any R.

For circuits that alternate with their time reverse, Shaw and Terhal reduce any number of rounds to a single round [9]; the gap argument above needs no such symmetry.

Circuits B and C. For circuit C, rounds are separated in time; a fault-free round inside the span of F splits it as in (2), so F occupies at most 11 consecutive rounds, and periodicity reduces any R to $R \le 13$. For circuit B we use a slab version of (2): a fault-free slab from the earliest reset to the latest measurement of one round splits F, provided that on every data qubit every propagating CNOT of round k inside the overlap with round $k + 1$ precedes every reading CNOT of round $k + 1$. We verified this condition for all rounds of the circuits with $R = 3, 8, 24$ in both bases; it reduces any R to $R \le 23$.

S8. UPPER BOUND

For each basis we found a weight-12 logical operator of the gross code of the corresponding type with CP-SAT and applied it as single-qubit flips immediately after data initialisation; each flip is one SD6 mechanism. For $R = 1, 3, 6, 10$ and both bases no detector fires and the logical observable flips; omitting any one of the twelve flips makes the error detectable. Hence $d_{\mathrm{circ}} \le 12$ for circuits A, B and C.

S9. FIG. 12 OF SBB: $d_{\mathrm{circ}} = 11$ AND RETIMINGS

$d_{\mathrm{circ}} = 11$. Fig. 12 is non-interleaved, so SBB's Prop. 3 and Cor. 1 give $d_{\mathrm{circ}} = d_{\mathrm{ext}}$ for every R; $d_{\mathrm{ext}} = 11$ by S4. The weight-11 extended logicals are realised by eleven circuit faults (checked by injection for $R = 1, 2, 3$ in both bases).

Lift structure. All 4224 weight-11 extended logicals per basis lift to undetected circuit errors of Fig. 12 ($R = 1, 2, 3$).

Retimings with period ≤ 8. Consider periodic schedules with period $P \leq 8$ shared by all checks, one reset, six CNOTs at distinct offsets and one measurement per check and period, and no data qubit in two CNOTs at one step. Weight 6 forces $P = 8$ and all offsets; keeping the Fig. 12 orders leaves one start per check. A backtracking search without symmetry assumptions finds exactly 32 conflict-free assignments, 4 up to global translation (two SAT encodings agree); all four are colour-uniform, deterministic and non-interleaved, and one is Fig. 12 itself. For each of them and every $R = 1, \ldots, 6$ in both bases we injected a weight-11 error that fires no detector and flips the observable.

Longer periods (bulk). For longer periods we use a Pauli-frame model of the bulk. A data X error at time θ on qubit w is first seen by Z check c (readouts at $T_c + nP$) in instance $H_c = \min\{n : T_c + nP > \theta\}$, i.e. for $\theta \in (T_c + (H_c - 1)P, T_c + H_c P)$. A weight-11 logical is *model-liftable* if instance indices H_c can be chosen such that the errors seen by each check can be paired, both errors of a pair being seen in the same instance (so all comparison detectors cancel), and, for every data error, the intervals of its three checks have a common point. For three open intervals of length P on the line, pairwise intersection, which is a difference constraint $H_c - H_d \leq \lfloor (T_d - T_c)/P \rfloor + 1$, implies a common point (one-dimensional Helly); an SD6 fault on w is available there (the two-qubit channel after the CNOT with the latest left endpoint contains $X_w \otimes I$). A residual of length 2 is deposited by an ancilla fault after the fourth CNOT of its check, which fixes its arrival times. *Direction used.* If the difference constraints of a logical of ten data errors and one length-2 residual are feasible, the corresponding eleven faults, placed away from the time boundaries, form a genuine undetected logical error; this was derived independently twice and confirmed by 75 direct injections, each firing no detector and flipping the observable. Non-existence statements need exactly this direction: an unsatisfiable search for a schedule in which no logical is model-liftable shows that every schedule has a model-liftable weight-11 logical, and when the search uses only logicals of the proven form, every schedule has a genuine weight-11 undetected error for sufficiently many rounds. The converse (model blocking implies exact blocking) is not used and fails at the initial time boundary, where the exact lift can succeed via first-round detectors. Blocking a model-liftable logical requires a negative cycle in its constraint graph; we encode this for all modelled logicals, for CP-SAT and in a lazy-clause SAT loop. The certificates block every model-liftable logical. Re-running the same encoders with only the 360 logicals per basis of the form ten data columns plus one length-2 residual gives UNSAT for every row of Table S3, each accompanied by a DRAT proof, so each obstruction rests only on the direction proven above. The resulting status for retimings that keep the Fig. 12 orders is given in Table S3. Per-check-free pipelined schedules with $P \geq 9$ and different CNOT orders are not covered; the latter is how the new circuits escape.

TABLE S3. Retimings of Fig. 12 that keep its CNOT orders. "Obstructed": no legal deterministic circuit of the family blocks every model-liftable weight-11 logical. Each row was also re-solved with only the ten-data-plus-one-length-2-residual logicals, which again gives UNSAT; every obstruction is therefore proven.

family	status
24-translation symmetric, common frame, any P	obstructed
24-translation symmetric, pipelined, $P = 8, \ldots, 12$	obstructed
subgroup-invariant pipelined, $P = 9$: three subgroups of order 12, three of order 6, one of order 8, two of order 4 and one of order 3	obstructed
subgroup-invariant pipelined, $P = 10$: one subgroup each of order 12, 8, 6	obstructed
no symmetry, common frame, $P \leq 16$	obstructed (one encoding)
other pipelined families with $P \geq 9$; symmetric pipelined $P \geq 13$	open

S10. SEARCH AND OPTIMALITY

Search. The circuits were found with SAT models over schedules symmetric under the 24 colour-preserving translations. Variables: for each of the six classes, one of the eight orders of the Fig. 12 class, a start and six CNOT times; constraints: legality (six distinct times modulo P per data qubit), determinism (on every overlapping pair of X and

Z checks, the number of shared qubits on which the X instance acts before the Z instance is even, for every instance offset), and, added lazily, one clause per frame-model-liftable weight-11 extended logical forbidding the timing that makes it liftable. Models were then certified with S5–S7. With the Fig. 12 orders pinned the model is unsatisfiable, in line with S9. As a control, 233 random legal deterministic schedules of the common-frame family, sampled without the blocking clauses, all have liftable weight-11 logicals.

Minimum period. With the same model and no circuit fixed, pipelined schedules with $P = 8, \ldots, 17$ and common-frame schedules with $P = 20$ are unsatisfiable. Every unsatisfiable clause set in the data set (129 in total: 104 for the obstruction and optimality instances with all logicals, and 25 length-2-only instances for Table S3) is accompanied by a DRAT proof produced by CaDiCaL and verified with drat-trim; the LRAT core was additionally verified with cake_lpr (SM S12). For the pipelined case a second, independent encoding (absolute times, one-sided difference literals) agrees for $P = 8, \ldots, 17$. The final clause sets were re-solved from scratch by independent SAT solvers (three for $P = 20$ common frame and $P = 12$, two for $P = 17$). Completeness of the family encoding: at $P = 8$ it enumerates exactly the 944 legal deterministic circuits found by an independent backtracking search, and 75 independently sampled legal deterministic circuits at $P = 12$, 17 (pipelined) and 20 (common frame) are all models. Soundness of the clauses: in 991 sampled cases, including logicals with several residual columns, every lift asserted by the frame model was confirmed on the exact DEM, and in 99 sampled cases every violated clause corresponded to an exact lift. Unsatisfiability means that every legal deterministic circuit of the family has a model-liftable weight-11 logical. For logicals of ten data errors and one length-2 residual this implies a genuine weight-11 undetected error for sufficiently many rounds (S9); for logicals with several residual columns the same implication is confirmed in the sampled cases above but not proven in general. The minimum periods are therefore established within the model, and are strong evidence for the circuits themselves; they are not statements at fixed small R. Restricting the CNOT window of class X_0 or of class Z_2 to at most 14 steps is unsatisfiable at $P = 24$, while each of the other four classes can be so restricted; since a window spans at most $P - 3$ steps, this suggests why $P \geq 18$ is needed, but the bound itself rests on the direct runs. At $P = 8$ all 944 legal deterministic circuits of the family were also checked directly: each has a weight-11 undetected error at $R = 1$ in the Z basis.

Uniqueness of the order class. Per check, the extended-code columns are the data columns and three residual sets; up to reversal there are 90 types per basis, of which 45 contain a residual of length 3 that, together with nine data columns, forms a weight-10 extended logical. For each of the 24 colourings by cosets of translation subgroups of index 2, 3, 4, 6, 8 and 9 (3, 4, 3, 12, 1 and 1 subgroups), with one type per colour, a SAT search with witness clauses leaves only the translates of Fig. 12 with $d_{\text{ext}} \geq 11$ in both bases; every witness used was re-verified as at most ten distinct columns of that assignment's extended code with vanishing H part and nonzero L part. Colourings of index 12, 18, 24 and 36 and per-check-free assignments are not covered.

Less symmetry. In the same model, pipelined schedules invariant under the order-12 subgroup generated by $(1, 2)$ are obstructed for every $P \leq 11$, and schedules invariant under any subgroup of order ≥ 2 are obstructed at $P = 8$. Schedules without symmetry, and the remaining subgroups at $P = 12, \ldots, 17$, are open.

S11. SENSITIVITY TO IDLE NOISE

Noise model. Starting from SD6 we multiply only the idle depolarizing channel by λ, $p_{\text{idle}} = \lambda p$; $\lambda = 1$ reproduces SD6 exactly and $\lambda = 0$ coincides with circuit-level depolarizing noise without idle errors (both checked on the detector error models). The CX-native variant uses SI1000 coefficients with CX, RX and MX treated as native operations; it is not the standard SI1000 model. Compiling CX to $H\,\mathrm{CZ}\,H$ and using Z-basis resets and measurements would raise the per-round fault mass of the period-18 circuit by about 12% and that of the period-8 circuits by 6–9%, so the variant slightly favours the period-18 circuit.

Direct sampling. $R = 12$, identical BP-OSD and Relay-BP decoders with exact priors, both bases. Table S4 lists per-round ratios of the logical error rate of the period-18 circuit to Fig. 12 and to the IBM circuit, with 95% intervals. At $\lambda = 0$ and $\lambda = 0.1$, $p = 4 \times 10^{-3}$, an independent re-sampling with a different noise implementation and BP-OSD code gave 0.37 [0.28, 0.49] and 2.73 [2.14, 3.47] for the ratio to Fig. 12. Several entries rest on few failures; at $p = 2 \times 10^{-3}$ and $\lambda = 0$ no circuit can be distinguished.

Per-layer rates. Dividing per-round rates by the period treats every layer as having unit duration, consistent with SD6 but not with physical time. Each round of the period-18 circuit has 12 CX-only layers and 6 layers with measurements or resets; Fig. 12 has 2 and 6. If measurement and reset layers are much longer than CX layers, the ratio of round durations falls from 2.25 towards 1.1–1.3, and the per-layer advantage of the longer period largely disappears.

Leading order. For an ideal decoder at low p, the per-round rate is dominated by weight-12 logical errors with total

TABLE S4. Per-round logical error rate of the period-18 circuit divided by that of Fig. 12 and of the IBM circuit (95% intervals; "–": too few failures).

noise	p	to Fig. 12, BP-OSD	to Fig. 12, Relay-BP	to IBM, BP-OSD	to IBM, Relay-BP
$\lambda = 0$	4×10^{-3}	0.34 [0.19, 0.59]	0.25 [0.11, 0.60]	0.89 [0.45, 1.7]	1.0 [0.33, 2.9]
$\lambda = 0$	3×10^{-3}	0.40 [0.16, 1.02]	0.40 [0.09, 1.9]	1.0 [0.34, 3.0]	0.50 [0.11, 2.4]
$\lambda = 0.1$	4×10^{-3}	3.0 [1.9, 4.7]	2.4 [1.3, 4.5]	3.7 [2.2, 6.0]	7.8 [2.9, 20]
$\lambda = 0.1$	3×10^{-3}	2.5 [0.81, 7.5]	4 [0.56, 22]	3.3 [0.96, 11]	–
$\lambda = 0.1$	2×10^{-3}	2.3 [1.09, 4.95]	–	1.1 [0.19, 4.9]	–
$\lambda = 0.3$	4×10^{-3}	33 [22, 50]	73 [42, 124]	44 [28, 68]	164 [79, 339]
$\lambda = 0.3$	3×10^{-3}	15 [6.9, 33]	20 [7.6, 51]	13 [6.3, 27]	–
CX-native	2×10^{-3}	23 [13, 39]	40 [22, 74]	36 [19, 70]	44 [23, 83]

weight $A_6 p^6$. Summing over the enumerated families (weight-11 logicals and one type of weight-12 logicals), the ratio $A_6(\text{period } 18)/A_6(\text{Fig. } 12)$ is 0.0104 at $\lambda = 0$, 0.084 at $\lambda = 0.1$, 0.84 at $\lambda = 0.3$ and 26 at $\lambda = 1$, crossing 1 near $\lambda \approx 0.32$. Caveats: A_6 is a union-type estimate, not a bound; the other type of weight-12 logical is not enumerated for either circuit, and an unknown factor F in their relative weight moves the crossing (to 0.50 for $F = 0.3$, 0.24 for $F = 2$, 0.20 for $F = 3$); the estimate describes $p \lesssim 10^{-4}$, not the sampled range (at $\lambda = 0.3$ it predicts near equality while direct sampling at 3–4×10^{-3} shows a factor 15–73); and it says nothing about practical decoders.

S12. VERIFICATION AND DATA

All circuits were built and simulated with Stim [15]; exact feasibility questions were decided with CP-SAT and with CDCL SAT solvers, and the extended-code searches use a dedicated depth-first search. The following were obtained with two independent implementations that agree: the construction of the extended code and $d_{\text{ext}} = 11$; the list of 4224 weight-11 logicals; the lift computation of Lemma 2; the enumeration of period-8 schedules; and the minimum-period runs. The arguments of Lemmas 1 and 3 were checked independently, and their structural premises (template identity, image set, periodicity, detector structure) were verified by computer for every circuit. The data set (Data availability) contains the circuit specifications (circuits A, B, C and Fig. 12), the lists of 4224 weight-11 extended logicals for each basis, the witnesses (the weight-12 upper-bound witnesses and the weight-11 witnesses for Fig. 12 and its retimings), the certificates (clause sets of the unsatisfiable instances), and the reproduction scripts for the lower-bound computations; it is available at [Zenodo DOI — PI to decide]. Every unsatisfiable clause set in the data set (129 in total: 104 for the obstruction and optimality instances with all logicals, and 25 length-2-only instances for Table S3), and every batch of the lift computation, is accompanied by a DRAT proof produced by CaDiCaL and verified with drat-trim; the LRAT core was additionally verified with cake_lpr. Here "accompanied" means: proofs of up to 100 MB are included in the data set; larger ones are identified by their SHA-256 hash and can be regenerated byte-identically with the included script. The certificates cover the solver decisions only; the circuit template and the reduction of Lemma 1 are covered by the pen-and-paper proofs and code-level checks of S5. Limitation: $d_{\text{ext}} = 11$ and the completeness of the list of 4224 weight-11 logicals rest on two independent exhaustive searches; we did not obtain a checkable proof certificate for them (a direct CNF encoding is beyond CDCL solvers already at weight 6).